

# Real-Time Sensor Analytics Dashboard – Technical Documentation

## 1. Introduction

This document describes the design and implementation of the **Proof of Concept (POC)** for a real-time sensor analytics dashboard.

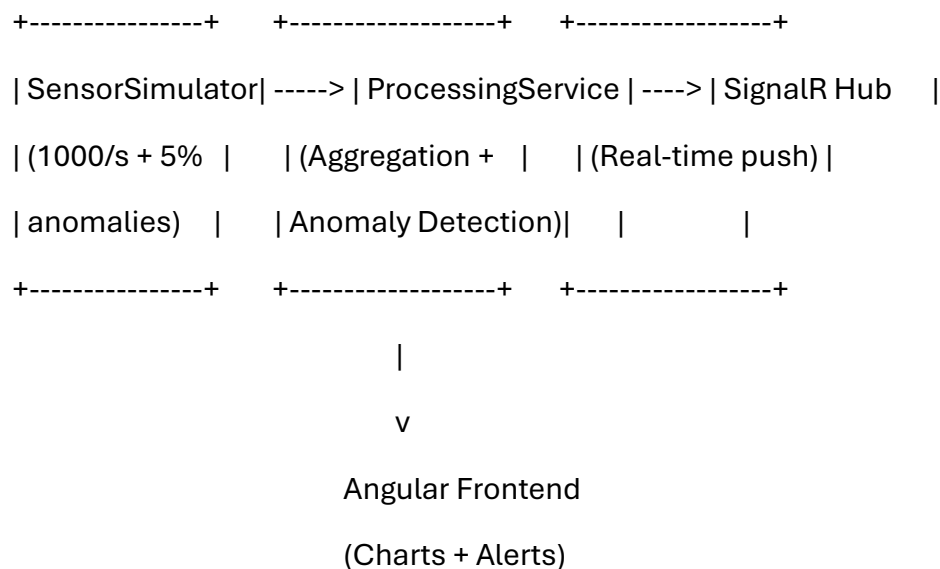
The objective is to:

- Ingest **1000 sensor readings per second**
- Process and detect anomalies in real time
- Retain up to **100,000 readings in memory**
- Stream results to an **Angular dashboard**
- Automatically purge readings older than 24 hours

---

## 2. System Architecture

### High-Level Diagram



---

## 3. Backend (.NET 9)

### Core Components

- **ASP.NET Core Web API** → Entry point for REST and SignalR

- **SignalR Hub (SensorHub)** → Real-time push to frontend
- **Channel<SensorReading>** → Internal high-throughput queue
- **CircularBuffer<SensorReading>** → Retains latest 100,000 readings in memory
- **Hosted Services:**
  - SensorSimulatorService → Produces synthetic sensor data
  - ProcessingService → Consumes, aggregates, detects anomalies, broadcasts
  - DataRetentionService → Purges data older than 24h (future extension)

## Data Flow

### 1. Simulation:

- Produces 1000 readings/sec across 3 sensors.
- 5% of readings are artificially spiked (to simulate anomalies).
- Each reading written into a Channel<SensorReading>.

### 2. Processing:

- ProcessingService reads continuously from the channel.
- Stores readings into a circular buffer (max 100k in memory).
- Aggregates batches of 100 readings:
  - avg, min, max, count, lastTimestamp per sensor.
- Passes values into the **AnomalyDetector**.

### 3. Anomaly Detection:

- Uses Z-score thresholding (simplified).
- Any value deviating strongly from recent mean triggers an anomaly.
- Broadcast anomaly event via SignalR.

### 4. Broadcasting:

- Aggregated batches sent as batch events.
  - Anomalies sent as anomaly events.
-

## 4. Frontend (Angular 19)

### Core Libraries

- **Angular 19 (standalone components)**
- **ng-apexcharts** → real-time time-series charts
- **PrimeNG** → UI components (cards, panels)
- **SignalR client** → for subscribing to backend streams

### UI Layout

- **Header:** Title + filters (time range, sensor selection)
- **KPI Cards:** Avg / Min / Max / Readings/sec per sensor
- **Charts:** One chart per sensor (scrolling, real-time)
- **Alerts Panel:** Shows anomalies in chronological order

### Data Flow in UI

- On load, connect to SignalR hub (/sensorhub).
  - Subscribe to:
    - "batch" → Update KPI cards + charts.
    - "anomaly" → Append entry in Alerts panel.
  - Charts update smoothly using **buffering and throttling** (to avoid UI lag).
- 

## 5. Data Model

### SensorReading

```
class SensorReading {  
    Guid SensorId { get; set; }  
    DateTime Timestamp { get; set; }  
    double Value { get; set; }  
    string Metadata { get; set; }  
}
```

### Aggregated Snapshot

```
{  
  "sensorId": "1111-1111...",  
  "avg": 20.3,  
  "min": 18.5,  
  "max": 22.0,  
  "count": 100,  
  "lastTimestamp": "2025-09-23T12:45:30Z"  
}
```

### Anomaly

```
{  
  "sensorId": "2222-2222...",  
  "value": 135.2,  
  "timestamp": "2025-09-23T12:45:32Z",  
  "type": "Critical"  
}
```

---

## 6. Performance Considerations

- **Throughput:**
  - Channel + async consumers → handles 1000/s easily.
- **Memory:**
  - Circular buffer ensures constant memory (~100k).
- **Batching:**
  - 100 readings grouped every 100ms → reduces frontend load.
- **Streaming:**
  - SignalR WebSocket → low latency push to UI.

---

## 7. Trade-offs

- **No external message queue (Kafka/RabbitMQ):**
  - Channels are enough for single-node POC.
  - For production, replace with Kafka for durability and scaling.
- **Anomaly detection is simple:**
  - Just threshold-based.
  - Real deployments may use ML/AI or adaptive thresholds.
- **100k buffer:**
  - Enough for demo and in-memory processing.
  - In real systems, would persist to a time-series DB (PostgreSQL, InfluxDB).

---

## 8. How to Extend

- **Persistence:** Store data in PostgreSQL (partitioned by timestamp).
- **Scaling:** Use Redis backplane for SignalR across multiple servers.
- **Alerts:** Add severity levels, cooldown, and historical audit.
- **Visualization:** Add multiple dashboards, filters, export features.

---

## 9. Validation

- Verified with **simulated load test**:
  - Sustains 1000 readings/sec.
  - UI updates without lag.
  - Anomalies (5% injection) appear correctly in Alerts panel.
- Stress-tested buffer with 200k events → old data drops as expected.

---

## 10. Conclusion

This POC demonstrates:

- **End-to-end streaming** from simulated sensors to a live Angular dashboard.
- Ability to handle high-frequency data (1000/s) in-memory.
- Real-time anomaly detection and visualization.
- Extensible design for persistence and scaling in production.